

# Software Defined Networking

## Lab Work 3 Introduction



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

Jeremias Blending, Leonhard Nobach, Christian Koch,  
Julius Rückert, Matthias Wichtlhuber



PS - Peer-to-Peer Systems Engineering Lab  
Dept. of Electrical Engineering and Information Technology  
Technische Universität Darmstadt  
Rundeturmstr. 12, D-64283 Darmstadt, Germany  
<http://www.ps.tu-darmstadt.de/>

# Lab 2 Task 1 Solution (1)

## ❖ Some confusion/uncertainty regarding the task

- We will stick to this line in the description: "A ping from h1 to h2 should result in a ping from h1 to h2 and h3. As a result, h1 receives more packets than it has sent."

```
mininet> h1 ping h2
64 bytes from 10.0.0.2: icmp_seq=343 ttl=64 time=0.076 ms
64 bytes from 10.0.0.2: icmp_seq=343 ttl=64 time=0.079 ms (DUP!)
```

- Verify using tcpdump

```
:~$ sudo tcpdump -ei s1-eth1
52 00:00:00:00:00:01 > 00:00:00:00:00:02, ethertype IPv4 (0x0800), length 98: 10.0.0.1 > 10.0.0.2:
ICMP echo request, id 5168, seq 48, length 64
92 00:00:00:00:00:02 > 00:00:00:00:00:01, ethertype IPv4 (0x0800), length 98: 10.0.0.2 > 10.0.0.1:
ICMP echo reply, id 5168, seq 48, length 64
95 00:00:00:00:00:03 > 00:00:00:00:00:01, ethertype IPv4 (0x0800), length 98: 10.0.0.3 > 10.0.0.1:
ICMP echo reply, id 5168, seq 48, length 64
```

# Lab 2 Task 1 Solution (2)



## ❖ We will accept all solutions that achieve that

- Ideal solution: install OpenFlow rules
- Other approaches
  - Duplicate packets on controller

```
def _duplicate_pingreq_to_host(pkg, target):  
    [...]  
    new_packet.add_protocol(ethernet.ethernet(ethertype=eth.ethertype,  
                                              dst=target[1],  
                                              src=eth.src))  
    new_packet.add_protocol(ipv4.ipv4(dst=target[0],  
                                       src=ipv4_frame.src,  
                                       proto=ipv4_frame.proto))  
    new_packet.add_protocol icmp.icmp(type_=icmp.ICMP_ECHO_REQUEST,  
                                       code=icmp_frame.code,  
                                       csum=icmp_frame.csum,  
                                       data=icmp_frame.data))  
    return new_packet
```

## ➤ Port mirroring, but no duplicate ping

```
[Somewhere inside add_flow]  
# dl_type required in order to use l3 protocol matching  
match_icmp = datapath.ofproto_parser.OFPMatch(  
    in_port=in_port, dl_type=0x0800, nw_proto=1)  
  
actions_icmp = [datapath.ofproto_parser.OFPACTIONOutput(ofproto.OFPP_FLOOD)]
```

# Lab 2 Task 1 Sample Solution (1)

- ❖ See the file „simple\_switch\_duplicate.py“ for a complete example solution
- ❖ The network is static and known: Add a host list

```
# Code by Patrick Welzel and Mahshid Okhovatzadeh
self.hosts_for_omniping = [
    {'port': 1,
     'ip': '10.0.0.1',
     'mac': '00:00:00:00:00:1'},
    {'port': 2,
     'ip': '10.0.0.2',
     'mac': '00:00:00:00:00:2'},
    {'port': 3,
     'ip': '10.0.0.3',
     'mac': '00:00:00:00:00:3'},
]
```

# Lab 2 Task 1 Sample Solution (2)



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

## ❖ Add additional output rules

```
Adarsh Chikkaballapur Umashankar
Bhargava Narasipura
actions=[]
# install a flow to avoid packet_in next time
if out_port != ofproto.OFPP_FLOOD:
    for host in self.hosts_for_omniping:
        if msg.in_port != host['port']:
            if out_port != host['port']:
                actions.append(datapath.ofproto_parser.OFActionSetDlDst(haddr_to_bin(host['mac'])))
                actions.append(datapath.ofproto_parser.OFActionSetNwDst(ofctl.ipv4_to_int(host['ip'])))
                actions.append(datapath.ofproto_parser.OFActionOutput(host['port']))
            self.logger.info("actions is %s ", actions)
            self.add_flow(datapath, msg.in_port, dst, actions)
```

## ❖ OpenFlow rules

```
mininet@mininet-vm:~$ sudo ovs-ofctl dump-flows s1
NXST_FLOW reply (xid=0x4):
 cookie=0x0, duration=41.308s, table=0, n_packets=32, n_bytes=3136, idle_timeout=20, idle_age=9,
 icmp,in_port=2 actions=output:1,mod_dl_dst:00:00:00:00:00:03,mod_nw_dst:10.0.0.3,output:3
 cookie=0x0, duration=41.311s, table=0, n_packets=32, n_bytes=3136, idle_timeout=20, idle_age=9,
 icmp,in_port=1 actions=output:2,mod_dl_dst:00:00:00:00:00:03,mod_nw_dst:10.0.0.3,output:3
 cookie=0x0, duration=41.308s, table=0, n_packets=32, n_bytes=3136, idle_timeout=20, idle_age=9,
 icmp,in_port=3 actions=output:1,mod_dl_dst:00:00:00:00:00:02,mod_nw_dst:10.0.0.2,output:2
```

# Lab Work 3

---



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

## Exploring OpenFlow 1.3

# Upgrading Mininet

- ❖ The Mininet version on your VM does not support OpenFlow 1.3
- ❖ We have to upgrade to the latest version
- ❖ Reboot your existing Mininet VM and enter:

```
~$ cd mininet
~/mininet$ git pull
~/mininet$ git checkout 2.2.0b1
~/mininet$ sudo ./util/install.sh -n
```

- ❖ Verify your installation

```
~/mininet$ mn --version
2.2.0b1
~/mininet$
```

# Run a simple\_switch with OpenFlow 1.0



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

## ❖ Test with OpenFlow 1.0

- First terminal: ssh session 1:

```
~$ ryu-manager ryu.app.simple_switch
```

- Second terminal: ssh session 2

```
~$ mn --topo single,3 --mac --arp --switch ovsk \  
    --controller=remote,ip=127.0.0.1  
mininet> h1 ping h2
```

- Third terminal: ssh session 3

- Note the new tool: ovs-ofctl (google „man ovs-ofctl“ for details)

```
~$ sudo ovs-ofctl dump-flows s1
```

## ❖ Note the debug output in session 1

- Investigate the source code
- Find out the meaning of the log messages



# Run a simple\_switch with OpenFlow 1.3



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

## ❖ Test with OpenFlow 1.3

### ➤ First terminal: ssh session 1:

```
~$ ryu-manager ryu.app.simple_switch_13
```

### ➤ Second terminal: ssh session 2

```
~$ mn --topo single,3 --mac --arp \  
    --switch ovsk,protocols=OpenFlow13 \  
    --controller=remote,ip=127.0.0.1  
mininet> h1 ping h2
```

### ➤ Third terminal: ssh session 3

```
~$ sudo ovs-ofctl dump-flows s1  
2014-11-24T11:47:01Z|00001|vconn|WARN|unix:/var/run/openvswitch/s1.mgmt: version negotiation  
failed (we support version 0x01, peer supports version 0x04)  
ovs-ofctl: s1: failed to connect to socket (Broken pipe)
```

# Run a simple\_switch with OpenFlow 1.3



TECHNISCHE  
UNIVERSITÄT  
DARMSTADT

## ➤ Third terminal: ssh session 3

```
~$ sudo ovs-ofctl dump-flows s1
```

```
2014-11-24T11:47:01Z|00001|vconn|WARN|unix:/var/run/openvswitch/s1.mgmt: version negotiation failed (we support version 0x01, peer supports version 0x04)  
ovs-ofctl: s1: failed to connect to socket (Broken pipe)
```

## ➤ Version negotiation failed

- We (ovs-ofctl) support version 0x01 (OpenFlow 1.0)
- Peer (Open vSwitch) supports version 0x04 (OpenFlow 1.3)
- ➔ note the version notation

```
~$ sudo ovs-ofctl dump-flows -O Openflow13 s1
```

```
OFFST_FLOW reply (OF1.3) (xid=0x2):  
cookie=0x0, duration=13.922s, table=0, n_packets=13, n_bytes=1274, priority=1,in_port=2,d1_dst=00:00:00:00:00:01 actions=output:1  
cookie=0x0, duration=12.924s, table=0, n_packets=12, n_bytes=1176, priority=1,in_port=1,d1_dst=00:00:00:00:00:02 actions=output:2  
cookie=0x0, duration=14.541s, table=0, n_packets=6, n_bytes=588, priority=0 actions=CONTROLLER:65535
```

## ➤ Specify the OpenFlow version "-O OpenFlow1X"

# **Task 1**


- ❖ Look at the source code of the Ryu modules `simple_switch.py` and `simple_switch_13.py`.
  - Describe and explain the differences between the OpenFlow 1.0 and OpenFlow 1.3 version of `simple_switch`.

# Task 2 (1)

## 1. Have a look at the OpenFlow 1.0-based filtering switch `simple_switch_filter.py`.

1. The basic forwarding method is again layer 2 switching.
2. We increase security by allowing only one host per port
3. This rule does not apply to inter-switch links
4. Run the topology:

```
sudo mn --topo linear,n=2,k=2 --mac --arp --switch ovsk \  
      --controller=remote,ip=127.0.0.1  
mininet> h1s1 ping h2s2
```



## 2. Verify

1. Install and run MAC spoofing tool

```
sudo apt-get update && sudo apt-get install nmap  
mininet> h1s1 nmap --spoof-mac 00:00:00:00:00:33 10.0.0.2
```

2. You should see the following message in your controller output:

```
dropping spoofed packet on s1 src=00:00:00:00:00:33 dst=ff:ff:ff:ff:ff:ff in_port=1
```

## **Task 2 (2)**

3. Have a look at the OpenFlow 1.0 based program that implements the filtering switch described in the last slide: `simple_switch_filter.py`.
4. Create an OpenFlow 1.3 version called `simple_switch_filter_13.py`.
  1. Base your implementation on `simple_switch_13.py` and `simple_switch_filter.py`
  2. Use two flow tables for the OpenFlow 1.3 version
    1. Use the first flow table for matching input ports and source MAC address
    2. Use the second flow table for sending the packets out the correct port.
5. Describe and discuss the differences regarding the number of flow rules used
  1. Consider different topologies and number of hosts for the sake of discussion even though running them with the `simple_switch_filter.py` is not possible
  2. Mathematically describe an upper bound for the number of flow rules
    1. For the OpenFlow 1.0 version
    2. For the OpenFlow 1.3 version

## Task 2



1. Have a look at the OpenFlow 1.0 based sample solution for lab 2 `simple_switch_duplicate.py`.
2. Create an OpenFlow 1.3 version of `simple_switch_duplicate.py` based on `simple_switch_13.py`.
  1. Use one flow table for matching of ports
  2. Use a second flow table for duplication and setting destination MAC and IP address
3. Describe and discuss the differences regarding the number of flow rules used
  1. Assume a large number of connected hosts for the sake of discussion

## ❖ Looking for a Python IDE?

- Eclipse with PyDev plugin
  - <http://pydev.org/>
- PyCharm (free for Student for research purposes)  
<https://www.jetbrains.com/student/>

## ❖ Things to be aware of

- Ryu relies on a concurrent networking library called Eventlet (<http://eventlet.net>)
- Debugging Eventlet requires special support and does currently neither work with PyDev nor with PyCharm out of the box